

Introduction to traditional record linkage and blocking techniques

Brenda Betancourt

Duke University
Department of Statistical Science
bb222@stat.duke.edu

February 8, 2018

Approximate matching

- Performing exact matching is very rare in practice since data is rarely noise/error free.
- Matching that allows fields to only be similar rather than exact duplicates.
- Most large-scale applications employ some level of approximate match between fields.
- Techniques for direct matching include edit distance and hashing algorithms for string data.

Example: Iterative deterministic linkage

Step 1: two records must match on SSN and one of the following:

- First and last name.
- Last name, month of birth, and sex.
- First name, month of birth, and sex.

Step 2: If SSN is missing or does not match, two records must match on last name, first name, month of birth, sex, and one of the following:

- Seven to eight digits of the SSN.
- Two or more of the following: year of birth, day of birth, middle initial, or date of death.

Approximate String Matching

Best for short strings such as person names:

- **Levenshtein (edit) distance (1966)**: minimum number of substitutions required to transform one string into another
e.g. Ad**a**m vs Al**a**n has a distance $L = 2$, normalized as $1 - \frac{L}{\max Length} = 0.5$.
- **Jaro-Winkler (1990)**: The Jaro distance (1989) considers common characters and character transpositions. The JW similarity measure is:

$$JW(A, B) = J(A, B) + 0.1p(1 - J(A, B))$$

where p is the # of the first four characters that agree exactly e.g.
Ad**a**m vs Al**a**n: $p=1$, $J= 0.67$ and $JW=0.7$.

Example: RLdata500

```
library(RecordLinkage)
data(RLdata500)
```

| ## | | fname_c1 | lname_c1 | by | bm | bd |
|----|-----|-----------|----------|------|----|----|
| ## | 314 | RENATE | SCHUTE | 1940 | 12 | 29 |
| ## | 407 | RENATE | SCHULTE | 1940 | 12 | 29 |
| ## | 289 | CHRISTINE | PETERS | 1993 | 2 | 5 |
| ## | 399 | CHRISTINE | PETERS | 1993 | 2 | 6 |
| ## | 402 | CHRISTA | SCHWARZ | 1965 | 7 | 13 |
| ## | 462 | CHRISTAH | SCHWARZ | 1965 | 7 | 13 |

Example: RLdata500

```
# Levenshtein similarity
```

```
levenshteinSim("SCHUTE", "SCHULTE")
```

```
## [1] 0.8571429
```

```
levenshteinSim("CHRISTA", "CHRISTAH")
```

```
## [1] 0.875
```

```
# Jaro-Winkler similarity
```

```
jarowinkler(c("SCHUTE", "CHRISTA"), c("SCHULTE", "CHRISTAH"))
```

```
## [1] 0.9714286 0.9750000
```

Soundex algorithm

- Generates a code that represents the phonetic pronunciation of a word, helps identifying spelling variations of names.
- The Soundex code for a name consists of a letter followed by three numerical digits:
 - the letter is the first letter of the name,
 - the digits encode the remaining consonants.
- Consonants at a similar place of articulation share the same digit
 - e.g. the labial consonants B, F, P and V are each encoded as the number 1.

Example: Soundex algorithm

| ## | fname_c1 | lname_c1 | by | bm | bd |
|--------|-----------|----------|------|----|----|
| ## 314 | RENATE | SCHUTE | 1940 | 12 | 29 |
| ## 407 | RENATE | SCHULTE | 1940 | 12 | 29 |
| ## 289 | CHRISTINE | PETERS | 1993 | 2 | 5 |
| ## 399 | CHRISTINE | PETERS | 1993 | 2 | 6 |
| ## 402 | CHRISTA | SCHWARZ | 1965 | 7 | 13 |
| ## 462 | CHRISTAH | SCHWARZ | 1965 | 7 | 13 |

```
tail(soundex(dup_set$fname_c1))
```

```
## [1] "R530" "R530" "C623" "C623" "C623" "C623"
```

```
tail(soundex(dup_set$lname_c1))
```

```
## [1] "S300" "S430" "P362" "P362" "S620" "S620"
```


Example: Soundex algorithm

| ## | fname_c1 | lname_c1 | by | bm | bd |
|--------|----------|----------|------|----|----|
| ## 130 | MICHAEL | MEYER | 1988 | 1 | 31 |
| ## 147 | MICHAEL | MYER | 1988 | 1 | 31 |
| ## 217 | HORST | MEIER | 1977 | 6 | 6 |
| ## 248 | HORST | MEIER | 1972 | 6 | 6 |
| ## 34 | HEINZ | BOEHM | 1938 | 12 | 20 |
| ## 111 | HEINZ | BOEHMR | 1938 | 12 | 20 |

```
head(soundex(dup_set$lname_c1))
```

```
## [1] "M600" "M600" "M600" "M600" "B500" "B560"
```

Blocking: Motivation

- Naively matching two files or finding duplicates within a file requires comparing all pairs of records.
- Infeasible for large files even when the comparisons are computationally inexpensive.
- The number of record pairs grows quadratically with the size of the dataset
 - Two files with 5,000 records \rightarrow 25,000,000 comparisons!


What is blocking?

Technique to reduce the comparison space:

- Filter out dissimilar record pairs that are extremely unlikely to be matches.
 - Perform record linkage only within blocks
- Traditional blocking : compare record pairs that match on one or more keys.
 - Creates a partition of the data
- Record pairs that do not meet the blocking criteria are automatically classified as non-matches.

Example: Traditional blocking

All-to-all record comparisons (left) versus partitioning records into blocks by lastname initial and comparing records only within each partition (right).



Continuation: RLdata500

```
# Record pairs for comparison  
choose(500,2)
```

```
## [1] 124750
```

```
# Blocking by last name initial  
last_init <- substr(RLdata500[, "lname_c1"], 1, 1)  
head(last_init)
```

```
## [1] "M" "B" "H" "W" "K" "F"
```

```
# Number of blocks  
length(unique(last_init))
```

```
## [1] 20
```

Continuation: RLdata500

```
# Number of records per block  
tbl <- table(last_init)  
head(tbl)
```

```
## last_init  
##   A   B   D   E   F   G  
##   5 56   2   6 38 12
```

```
# Block sizes can vary a lot  
summary(as.numeric(tbl))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##      2.00   5.75    8.00   25.00   40.00   115.00
```

Continuation: RLdata500

```
# Number of records pairs per block  
sapply(tbl, choose, k=2)
```

| | | | | | | | | | | | |
|----|------|------|----|------|-----|----|-----|----|------|----|------|
| ## | A | B | D | E | F | G | H | J | K | L | M |
| ## | 10 | 1540 | 1 | 15 | 703 | 66 | 496 | 28 | 1035 | 78 | 2850 |
| ## | S | T | V | W | Z | | | | | | |
| ## | 6555 | 1 | 21 | 1326 | 10 | | | | | | |

```
# Reduction on comparison space  
sum(sapply(tbl, choose, k=2))
```

```
## [1] 14805
```

Blocking caveats

- Fields can be unreliable for many applications and blocking may miss large proportions of matches i.e. increased false negatives rates.
- The frequency distribution of the values in the fields used as blocking keys will affect the size of the blocks.
- Trade-off between block sizes: true matches being missed vs computational efficiency.

How to choose the blocking key or keys

- Fields containing the fewest errors or missing values should be chosen as blocking variables e.g. clinical diagnosis in EHR.
- Understand the kinds of errors that are unlikely for a certain field or a combination of them.
- More complex blocking schemes can be constructed using conjunctions.
 - Retain only pairs which agree on either last name initial and zip code

Example: Voter Survey data

The Views of the Electorate Research (VOTER) Survey was conducted by the survey firm YouGov.

- 8,000 adults (age 18+) with internet access took the survey on-line between November 29 and December 29, 2016.
- These respondents were originally interviewed by YouGov in 2011-2012.
- Barack Obama (Democrat) won in 2012 and Donald Trump (Republican) won in 2016.

Continuation: Voter Survey data

- Demographic variables
 - Year of birth (age)
 - Gender
 - Race
 - State
 - Education level
 - Family income
- Party affiliation: democrat, republican, independent, other

Which fields are reliable for blocking in this example?

Continuation: Is race reliable?

| | 2012 | 2016 |
|-----------------|------|------|
| White | 6244 | 6198 |
| Black | 654 | 645 |
| Hispanic | 400 | 397 |
| Mixed | 160 | 186 |
| Other | 137 | 167 |
| Asian | 117 | 118 |
| Native American | 60 | 59 |
| Middle Eastern | 10 | 12 |

| | White | Black | Mixed | Other |
|-------|-------|-------|-------|-------|
| White | 6073 | 5 | 46 | 74 |
| Black | 4 | 627 | 10 | 10 |
| Mixed | 31 | 6 | 100 | 8 |
| Other | 50 | 4 | 14 | 62 |

Continuation: Is party affiliation reliable?

| | Democrat | Indepen. | Republican | Not sure | Other |
|------------|----------|----------|------------|----------|-------|
| Democrat | 2424 | 192 | 90 | 25 | 23 |
| Indepen. | 263 | 1929 | 221 | 16 | 57 |
| Republican | 39 | 215 | 1881 | 11 | 60 |
| Not sure | 48 | 48 | 54 | 41 | 5 |
| Other | 17 | 46 | 34 | 2 | 41 |

Blocking by disjunctions

- Produces overlapping blocks of the data.
- Using multiple keys to consider typographical or measurement errors that would exclude true matches.
 - Blocking by last name initial or zip code

| | | | |
|----------|------------|----------------|-------|
| <i>A</i> | Mary Clain | 123 Oak St | 90210 |
| <i>B</i> | Mary Klein | 123 Oak Street | 90210 |
| <i>C</i> | Mary Klain | 123 Oak St | 50210 |

- Reduction in false negative rates.

Example: Blocking by disjunctions

```
# Two records must agree in either first name initial  
# or birth year to be compared.  
# Only 2709 pairs instead of 124750!
```

```
rpairs <- compare.dedup(RLdata500c,  
blockfld = list(1, 3), #list with blocking fields  
identity = identity.RLdata500)
```

```
tail(rpairs$pairs)
```

| | ## | id1 | id2 | fname_c1 | lname_c1 | by | bm | bd | is_match |
|----|------|-----|-----|----------|----------|----|----|----|----------|
| ## | 2704 | 477 | 497 | 1 | 0 | 0 | 0 | 0 | 0 |
| ## | 2705 | 479 | 483 | 0 | 0 | 1 | 1 | 0 | 0 |
| ## | 2706 | 480 | 481 | 1 | 0 | 0 | 0 | 0 | 0 |
| ## | 2707 | 480 | 490 | 1 | 0 | 0 | 0 | 0 | 0 |
| ## | 2708 | 481 | 490 | 1 | 1 | 0 | 1 | 1 | 1 |
| ## | 2709 | 494 | 497 | 0 | 0 | 1 | 1 | 0 | 0 |

Example: String comparison and blocking

```
rpairsfuzzy <- compare.dedup(RLdata500c, phonetic = FALSE,  
blockfld = 3, strcmp = TRUE, strcmpfun = jarowinkler)  
  
tail(rpairsfuzzy$pairs)
```

| ## | | id1 | id2 | fname_c1 | lname_c1 | by | bm | bd | is_match |
|----|------|-----|-----|-----------|-----------|----|-----|-----|----------|
| ## | 1540 | 460 | 485 | 0.0000000 | 0.5396825 | 1 | 0.7 | 0.0 | NA |
| ## | 1541 | 464 | 466 | 0.4555556 | 0.5396825 | 1 | 0.7 | 0.0 | NA |
| ## | 1542 | 467 | 472 | 1.0000000 | 0.9333333 | 1 | 1.0 | 1.0 | NA |
| ## | 1543 | 468 | 469 | 0.5777778 | 0.4666667 | 1 | 0.7 | 0.7 | NA |
| ## | 1544 | 479 | 483 | 0.4370370 | 0.5619048 | 1 | 1.0 | 0.0 | NA |
| ## | 1545 | 494 | 497 | 0.6111111 | 0.5026455 | 1 | 1.0 | 0.0 | NA |

The Fellegi-Sunter approach (1969)

- Represent every pair of records using vector of features that describe similarity between individual record fields.
 - Use string metrics (Jaro-Winkler) and edit-distances for names and strings of numbers.
- Place feature vectors for record pairs into three classes: matches (M), nonmatches (U), and possible matches.
- Let $P(\gamma|M)$ and $P(\gamma|U)$ be probabilities of observing a feature vector γ for a matched and nonmatched pair, respectively.

The Fellegi-Sunter approach (1969)

- Perform record-pair classification by calculating the ratio $w = (P(\gamma|M)/P(\gamma|U))$ for each candidate record pair.
- Establish two thresholds based on desired error levels to optimally separate the weight values for matches, possibly matches, and nonmatches.
- **Drawbacks:** only for two files, no transitive closures.

Example: Fellegi-Sunter

```
#tail(rpairs$pairs)  
# Using comparison data blocking by first name initial  
# and birth year  
rpairs1 <- epiWeights(rpairs)  
  
# Weights to compute thresholds for classification  
head(rpairs1$Wdata)
```

```
## [1] 0.2223402 0.2223402 0.2488181 0.2488181 0.3936336 0.
```

Example: Fellegi-Sunter

```
summary(rpairs1)
```

Weight distribution:

| | | | | |
|------------|------------|------------|------------|------------|
| (0.35,0.4] | (0.4,0.45] | (0.45,0.5] | (0.55,0.6] | (0.6,0.65] |
| 2 | 10 | 30 | 50 | 8 |
| (0.65,0.7] | (0.7,0.75] | (0.75,0.8] | (0.8,0.85] | (0.85,0.9] |
| 0 | 0 | 35 | 8 | 3 |

Example: Fellegi-Sunter

```
result <- epiClassify(rpairs1, 0.7)
summary(result)
```

```
alpha error: 0.080000 # False negative rate
beta error:  0.000000 # False positive rate
accuracy: 0.998523
```

Classification table:

| | | classification | | |
|-------------|------|----------------|----|---|
| true status | | N | P | L |
| FALSE | 2659 | 0 | 0 | |
| TRUE | 4 | 0 | 46 | |

References

- Sariyar M. and Borg A. (2010), The RecordLinkage Package: Detecting Errors in Data, The R Journal Vol. 2/2. **Check big data functions of the package ($\geq 1'000,000$ record pairs).**
- Steorts et al. (2014) A Comparison of Blocking Methods for Record Linkage. In: Domingo-Ferrer J. (eds) Privacy in Statistical Databases. PSD 2014. Lecture Notes in Computer Science, vol 8744. Springer, Cham.
- Dusetzina et al. Linking Data for Health Services Research: A Framework and Instructional Guide. Rockville (MD): Agency for Healthcare Research and Quality (US) (2014). An Overview of Record Linkage Methods.
<https://www.ncbi.nlm.nih.gov/books/NBK253312/>
- Entity Resolution and Information Quality, John R. Talburt, Elsevier (2011)